

# Replicating DreamBooth: Subject-Driven Fine-Tuning of Stable Diffusion

Ryan Qiu   Gordon Mei   Caleb Shim   Evan Cui

Cornell University   CS4782 — Deep Learning

<https://github.com/caleb-bit/dreambooth-replication>

---

## Introduction/Motivation

Large-scale text-to-image diffusion models synthesize photorealistic images from prompts but cannot render a *specific* subject unseen during training. Prompting “a photo of my dog in Paris” produces a plausible dog, but not the specific dog that the user wants. Even exhaustive description (breed, color, distinguishing markings) cannot uniquely identify a specific instance, since the model has no reference to bind to. **DreamBooth** [1] (Ruiz et al., CVPR 2023) addresses this by fine-tuning the model on 3–5 reference images, binding a particular subject to a rare identifier token and enabling compositional generation in novel contexts.

We replicate DreamBooth on Stable Diffusion v1.5 (SD 1.5) [2], implementing the full training pipeline from scratch. Our goal is to reproduce the paper’s reported quantitative metrics on SD 1.5 and document the implementation details necessary to achieve them.

## Chosen Result

We target the quantitative evaluation from Table 2 of the DreamBooth paper, which reports subject fidelity (DINO [4], CLIP-I [3]) and prompt fidelity (CLIP-T) for fine-tuned SD 1.5 models across 30 subjects. The paper reports mean scores of DINO 0.668, CLIP-I 0.803, and CLIP-T 0.305 for SD. These results show that a fine-tuned model can preserve specific subject identity while remaining responsive to novel prompts. Reproducing these scores provides evidence that our pipeline correctly implements key training components, such as rare-token binding and prior preservation mechanisms.

## Methodology

**Architecture.** SD 1.5 has three components: a VAE ( $\approx 83$ M params) that encodes images to latents, a UNet ( $\approx 860$ M params) that denoises in latent space, and a CLIP text encoder ( $\approx 123$ M params). We fine-tune the UNet only; both the VAE and CLIP text encoder are frozen throughout training.

**Training objective.** Instance images are captioned as “a photo of **sks** [class]”, where **sks** is a rare identifier with negligible prior semantics. The model has no strong associations with this token, so it can be used to represent the new subject. The class label (e.g. “dog”) anchors **sks** within the correct semantic area, ensuring the model understands the general category while learning to associate the rare token with the specific instance. At each step, both instance and class images are encoded to latents, noise is sampled and added, and the UNet predicts the noise. The combined loss is:

$$\mathcal{L} = \underbrace{\|\hat{\varepsilon}_{\text{inst}} - \varepsilon_{\text{inst}}\|^2}_{\text{instance loss}} + \lambda \underbrace{\|\hat{\varepsilon}_{\text{class}} - \varepsilon_{\text{class}}\|^2}_{\text{prior preservation}}, \quad \lambda = 1.0.$$


Prior-preservation class images (100 per subject) are generated by the unmodified SD 1.5 model before training, substituting for the Google Imagen model used in the original paper.

**Hyperparameters.** 800 steps, learning rate  $5 \times 10^{-6}$ , resolution  $512 \times 512$ , batch size 1, 8-bit AdamW optimizer [5], gradient checkpointing. The UNet runs in 32-bit floating point precision (FP32) while frozen components use 16-bit precision (FP16) for memory efficiency.

**Dataset and evaluation.** We fine-tune on **dog**, **dog2**, and **backpack** from the official DreamBooth dataset [7], each with 5 instance images. Evaluation generates 4 images per prompt across 5 prompts (20 images/subject). DINO uses ViT-S/16 [4]; CLIP-I and CLIP-T use CLIP ViT-L/14 [3]. Training runs on a Google Colab T4 GPU ( $\approx 35$  min/subject); metrics are computed on CPU.

## Results & Analysis

Our mean DINO (0.598) falls slightly below the paper’s 0.668, expected given that we evaluate only 3 subjects rather than the full 30. CLIP-I (0.861) exceeds the paper’s 0.803, indicating strong semantic alignment. CLIP-T (0.265) is below the paper’s 0.305, reflecting the subject–prompt fidelity trade-off. Qualitatively, generated images maintain subject appearance across novel backgrounds and lighting; see Figure 1. The most notable gap is DINO variance across subjects: **dog** scores 0.716 (near the paper’s mean) while **dog2** and **backpack** score around 0.54, suggesting fine-grained identity binding is sensitive to reference image consistency. The backpack’s high CLIP-I (0.887) despite lower DINO indicates semantic similarity is preserved even where exact feature identity is weaker.



(a) Reference subject



(b) “a photo of [v] dog with Cornell merch”

Figure 1: Reference image and a generated image using the fine-tuned model, demonstrating subject identity preserved in a novel context.

## Reflections

**Prior preservation.** Prior loss is a critical part of the loss function that ensures that the model maintains its class representations. Without the prior preservation loss, fine-tuning on the 3–5 reference images will cause the model to significantly overfit, where the rare token begins to displace the entire class representation (e.g. prompting for “a photo of a dog” returns the specific subject instead of an arbitrary dog).

**Using 16-bit floating point precision causes silent gradient underflow.** Using half-precision for the UNet caused gradient updates to fall below the representable range, halting weight updates with no error or warning. Keeping the UNet in 32-bit precision resolved this.

**Freezing the text encoder.** The text encoder’s job is to map language to a fixed semantic space, a structure that is already well-learned and should not change. Fine-tuning it allows the model to take a shortcut: rather than learning to visually bind the subject through the UNet, it just redefines what **sks** means at the language level, which destroys its ability to compose with novel prompts.

**Extension: Logo fine-tuning.** We applied DreamBooth to a custom CS4782 course logo using 5 augmented variants as instance images, with **sks** as the identifier and **logo** as the class noun. Direct

|                    | DINO↑        | CLIP-I↑      | CLIP-T↑      |
|--------------------|--------------|--------------|--------------|
| dog                | 0.716        | 0.867        | 0.275        |
| dog2               | 0.542        | 0.828        | 0.247        |
| backpack           | 0.536        | 0.887        | 0.274        |
| <b>Ours (mean)</b> | <b>0.598</b> | <b>0.861</b> | <b>0.265</b> |
| Real images        | 0.774        | 0.885        | —            |
| Paper (SD)         | 0.668        | 0.803        | 0.305        |

Table 1: Quantitative metrics. “Real images” is the upper bound; “Paper (SD)” is the DreamBooth paper’s SD 1.5 result.

generation (“a photo of a **sks** logo”) was able to reproduce the graphic faithfully, but not with complete accuracy. Context placement failed: prompts such as “a man wearing a **sks** logo” or “a mug with a **sks** logo” caused the model to render **sks** as literal glyphs rather than the learned graphic (see Figure 2). It may be useful to include contextual mockups in the reference set (e.g., the logo composited onto a shirt, paired with a matching prompt), so the model sees the subject in context during training rather than only as an isolated graphic.

**Style transfer.** We also applied Dreambooth to generate new paintings in the art style of Van Gogh, with **sks** as the identifier and **painting** as the class noun. One important lesson is the necessity of tweaking the hyperparameters, since a style is more abstract than an object. Our model for style transfer was trained for 600 steps with learning rate  $2 \times 10^{-6}$ , as training with the hyperparameters used for the dog and backpack models led to significant overfitting (i.e. every prompt for Van Gogh style landscape led to Starry Night being generated). Still, some of the visible lines in the Van Gogh-style backgrounds like the city paintings suggest the model is slightly overfit, as the model tries to force the generated skyline to match the subject images too hard. Despite this, the model qualitatively seemed to perform well at generating Van Gogh-style paintings, such as the beach, car and person paintings in the results folder. The model was also still able to generate paintings in styles other than Van Gogh, as seen in the paintings of a person generated without using the identifier. The key takeaway is that Dreambooth can be used effectively for style transfer. However, it is important to note that Dreambooth may not be the best option, as LoRA models for style transfer also exist that may be faster to train and/or produce better results, though we did not do much research into them.

**Future work.** We want to scale to 30 subjects, giving us a direct benchmark comparison against the paper’s reported results. In addition, we would like to try multi-subject fine-tuning, where we bind multiple rare tokens (e.g., **sks** and **xyz**) to distinct subjects in a single fine-tune. This would test whether the model can compose two learned identities simultaneously.




Figure 2: Logo fine-tuning results. **Top row:** reference logo followed by four samples from “a photo of a **sks** logo” — direct generation reproduces the graphic with mixed success. **Bottom row:** context placement prompts cause the model to render **sks** as literal glyphs rather than the learned logo.

## References

- [1] N. Ruiz, Y. Li, V. Jampani, Y. Pritch, M. Rubinstein, and K. Aberman. DreamBooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *CVPR*, 2023.
- [2] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. In *CVPR*, 2022.
- [3] A. Radford et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- [4] M. Caron et al. Emerging properties in self-supervised vision transformers. In *ICCV*, 2021.
- [5] T. Dettmers, M. Lewis, Y. Belkada, and L. Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *NeurIPS*, 2022.
- [6] HuggingFace. Diffusers: State-of-the-art diffusion models. <https://github.com/huggingface/diffusers>, 2022.
- [7] Google. DreamBooth dataset. <https://github.com/google/dreambooth>, 2022.