

Replicating DreamBooth: Subject-Driven Fine-Tuning of Stable Diffusion

Ryan Qiu Gordon Mei Caleb Shim Evan Cui
Cornell University CS4782 — Deep Learning

<https://github.com/caleb-bit/dreambooth-replication>

Introduction/Motivation

Large-scale text-to-image diffusion models synthesize photorealistic images from prompts but cannot render a *specific* subject unseen during training. Prompting “a photo of my dog in Paris” produces a plausible dog, but not the specific dog that the user wants. Even exhaustive description (breed, color, distinguishing markings) cannot uniquely identify a specific instance, since the model has no reference to bind to. **DreamBooth** [1] (Ruiz et al., CVPR 2023) addresses this by fine-tuning the model on 3–5 reference images, binding a particular subject to a rare identifier token and enabling compositional generation in novel contexts.

We replicate DreamBooth on Stable Diffusion v1.5 (SD 1.5) [2], implementing the full training pipeline from scratch. Our goal is to reproduce the paper’s reported quantitative metrics on SD 1.5 and document the implementation details necessary to achieve them.

Chosen Result

We target the quantitative evaluation from Table 2 of the DreamBooth paper, which reports subject fidelity (DINO [4], CLIP-I [3]) and prompt fidelity (CLIP-T) for fine-tuned SD 1.5 models across 30 subjects. The paper reports mean scores of DINO 0.668, CLIP-I 0.803, and CLIP-T 0.305 for SD. These results show that a fine-tuned model can preserve specific subject identity while remaining responsive to novel prompts. Reproducing these scores provides evidence that our pipeline correctly implements key training components, such as rare-token binding and prior preservation mechanisms.

Methodology

Architecture. SD 1.5 has three components: a VAE (≈ 83 M params) that encodes images to latents, a UNet (≈ 860 M params) that denoises in latent space, and a CLIP text encoder (≈ 123 M params). We fine-tune the UNet only; both the VAE and CLIP text encoder are frozen throughout training.

Training objective. Instance images are captioned as “a photo of **sks** [class]”, where **sks** is a rare identifier with negligible prior semantics. The model has no strong associations with this token, so it can be used to represent the new subject. The class label (e.g. “dog”) anchors **sks** within the correct semantic area, ensuring the model understands the general category while learning to associate the rare token with the specific instance. At each step, both instance and class images are encoded to latents, noise is sampled and added, and the UNet predicts the noise. The combined loss is:

$$\mathcal{L} = \underbrace{\|\hat{\varepsilon}_{\text{inst}} - \varepsilon_{\text{inst}}\|^2}_{\text{instance loss}} + \lambda \underbrace{\|\hat{\varepsilon}_{\text{class}} - \varepsilon_{\text{class}}\|^2}_{\text{prior preservation}}, \quad \lambda = 1.0.$$

Prior-preservation class images (100 per subject) are generated by the unmodified SD 1.5 model before training, substituting for the Google Imagen model used in the original paper.

Hyperparameters. 800 steps, learning rate 5×10^{-6} , resolution 512×512 , batch size 1, 8-bit AdamW optimizer [5], gradient checkpointing. The UNet runs in 32-bit floating point precision (FP32) while frozen components use 16-bit precision (FP16) for memory efficiency.

Dataset and evaluation. We fine-tune on **dog**, **dog2**, and **backpack** from the official DreamBooth dataset [7], each with 5 instance images. Evaluation generates 4 images per prompt across 5 prompts (20 images/subject). DINO uses ViT-S/16 [4]; CLIP-I and CLIP-T use CLIP ViT-L/14 [3]. Training runs on a Google Colab T4 GPU (≈ 35 min/subject); metrics are computed on CPU.

Table 1: Quantitative metrics. “Real images” is the upper bound; “Paper (SD)” is the DreamBooth paper’s SD 1.5 result.

	DINO↑	CLIP-I↑	CLIP-T↑
dog	0.716	0.867	0.275
dog2	0.542	0.828	0.247
backpack	0.536	0.887	0.274
Ours (mean)	0.598	0.861	0.265
Real images	0.774	0.885	—
Paper (SD)	0.668	0.803	0.305

Results & Analysis

Our mean DINO (0.598) falls slightly below the paper’s 0.668, which is expected given that we evaluate only 3 subjects rather than the full 30. CLIP-I (0.861) exceeds the paper’s 0.803, indicating strong semantic alignment between generated and reference images. CLIP-T (0.265) is below the paper’s 0.305, reflecting the inherent subject–prompt fidelity trade-off: our fine-tuned models preserve identity well but sacrifice some text adherence. Qualitatively, generated images maintain subject appearance across novel backgrounds and lighting; see Figure 1.

The most notable gap is DINO variance across subjects: **dog** scores 0.716 (near the paper’s mean) while **dog2** and **backpack** score around 0.54, suggesting that the model’s ability to maintain feature identity is highly dependent on the reference images. The backpack scores high on CLIP-I (0.887) despite lower DINO, indicating that semantic similarity is preserved even where exact feature identity is weaker.

Overall, this replication demonstrated that we are able to achieve significant evaluation thresholds for subject-driven generation even when training on a smaller diffusion model like SD 1.5 on a smaller set of generated samples.


Reflections

Prior preservation. Prior loss is a critical part of the loss function that ensures that the model maintains its class representations. Without the prior preservation loss, fine-tuning on the 3–5 reference images will cause the model to significantly overfit, where the rare token begins to displace the entire class representation (e.g. prompting for “a photo of a dog” returns the specific subject instead of an arbitrary dog).

Using 16-bit floating point precision causes silent gradient underflow. Using half-precision for the UNet caused gradient updates to fall below the representable range, halting weight updates with no error or warning. Keeping the UNet in 32-bit precision resolved this.

Freezing the text encoder. The text encoder’s job is to map language to a fixed semantic space, a structure that is already well-learned and should not change. Fine-tuning it allows the model to take a shortcut: rather than learning to visually bind the subject through the UNet, it just redefines what **sks** means at the language level, which destroys its ability to compose with novel prompts.

Extension: Logo fine-tuning. We applied DreamBooth to a custom CS4782 course logo using 5 augmented variants as instance images, with **sks** as the identifier and **logo** as the class noun. Direct generation (“a photo of a **sks** logo”) was able to reproduce the graphic faithfully, but not with complete accuracy. Context placement failed: prompts such as “a man wearing a **sks** logo” or “a mug with a **sks** logo” caused the model to render **sks** as literal glyphs rather than the learned graphic (see Figure 2). It may be useful to include contextual mockups in the reference set (e.g., the logo composited onto a shirt, paired with a matching prompt), so the model sees the subject in context during training rather than only as an isolated graphic.



(a) Reference subject




(b) “a photo of [v] dog with Cornell merch”

Figure 1: Reference image and a generated image using the fine-tuned model, demonstrating subject identity preserved in a novel context.



(a) Reference




(b)



(c)



(d)



(e)



(f) “a man wearing a sks logo”



(g) “a mug with a sks logo”

Figure 2: Logo fine-tuning results. **Top row:** reference logo followed by four samples from “a photo of a **sks** logo” — direct generation reproduces the graphic with mixed success. **Bottom row:** context placement prompts cause the model to render **sks** as literal glyphs rather than the learned logo.

References

- [1] N. Ruiz, Y. Li, V. Jampani, Y. Pritch, M. Rubinstein, and K. Aberman. DreamBooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *CVPR*, 2023.
- [2] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. In *CVPR*, 2022.
- [3] A. Radford et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- [4] M. Caron et al. Emerging properties in self-supervised vision transformers. In *ICCV*, 2021.
- [5] T. Dettmers, M. Lewis, Y. Belkada, and L. Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *NeurIPS*, 2022.
- [6] HuggingFace. Diffusers: State-of-the-art diffusion models. <https://github.com/huggingface/diffusers>, 2022.
- [7] Google. DreamBooth dataset. <https://github.com/google/dreambooth>, 2022.